Response Letter with Annotations

Citation (APA format, using Zotero or other bibliography software)

King, G., Pan, J., & Roberts, M. E. (2013). How Censorship in China Allows Government Criticism but Silences Collective Expression. *American Political Science Review*, *107*(2), 326–343. https://doi.org/10.1017/S0003055413000014

Annotations

1. **Purposes**
    a. Motivation (one-sentence summary)

The motivation of this study was to identify the underlying logic of Chinese government censorship by analyzing which types of social media content are systematically removed using large-scale automated data collection.

b. Specific Rationale (several bullet points)

- Censorship mechanism is not transparent in China
- Censorship requires manual labor, therefore computation methods may be able to detect certain topics that are more likely to be censorship with greater speed.
- To test two theories "state critique theory" and "theory of collective action potential"

c. Research Questions and/or Hypotheses

The central hypothesis is that the government censors all posts in topic areas during volume bursts that discuss events with collective action potential.
Authors hypothesize that the censors make the much easier judgment, about whether the posts are on topics associated with events that have collective action potential, as they do it regardless of whether or not the posts criticize the state.

- Identify IV, DV, mediator, moderator

IV - five content areas: Collective action potential, criticism of the censors, pornography, government policies and other news.
DV - posts removed
(Note: This study is largely descriptive and focuses on counts and comparisons of posts that were removed.)

- Identify specific relationships among variables

2. **Method**
    a. Research Design
        - Experiment: conditions, manipulations, etc.
        - Survey: cross-sectional, longitudinal, etc.
        - Content analysis: data source, etc.
        - …

The study uses a computational, observational design that tracks social media posts over time to compare censored and non-censored content.

Data collection method was able to outpace the manual censorship process.

b. Sample (representative sample?)

Accessed 1,382 Chinese websites. Stratified random sampling design, organized hierarchically. Collected over 3.6 million posts and randomly selected 127,283 for analysis. Reasonably representative.

c. Measures

Posts - removed or not
Censorship magnitudes - Bursts
Topics - 85 areas
Political sensitivity: High, medium and low.
Content areas - Topics classified into: Collective action potential, criticism of the censors, pornography, government policies and other news.
Collective action potential: involve protest or organized crowd formation outside the internet, relate to individuals who have organized or incited collective action on the ground in the past, or relate to nationalism or nationalist sentiment that have incited protest or collective action in the past.

d. Analysis

- Calculated overall censorship rates across all posts to establish baseline levels of removal.
- Identified volume bursts using time-series outlier detection to locate periods of unusually high posting activity.
- Compared censorship rates inside versus outside volume bursts.
- Computed censorship magnitude as the difference in percentage of posts censored during bursts versus normal periods.
- Compared censorship patterns across event categories (collective action, government policy, news, etc.).
- Applied automated text analysis (Hopkins & King method) to classify posts as critical, supportive, or neutral toward the state.
- Compared censorship rates between critical and supportive posts.

-

3. Results and Conclusions

a. Answers to RQs and Hypotheses

- Consistent with the theory of collective action potential, some of the most highly censored events are not criticisms or even discussions of national policies, but rather highly localized collective expressions that represent or threaten group formation.
- In all events categorized as having collective action potential, censorship within the event is more frequent than censorship outside the event.
- According to the all four graphs presented, the pattern indicates that the Chinese authorities disproportionately focus considerable censorship efforts during volume bursts.
- The level of censorship for two topics (pornography and criticism of the censors) remained high consistently during the data collection period.

b. Include Key Figures/Tables

Figure 1. The Fractured Structure of the Chinese Social Media Landscape

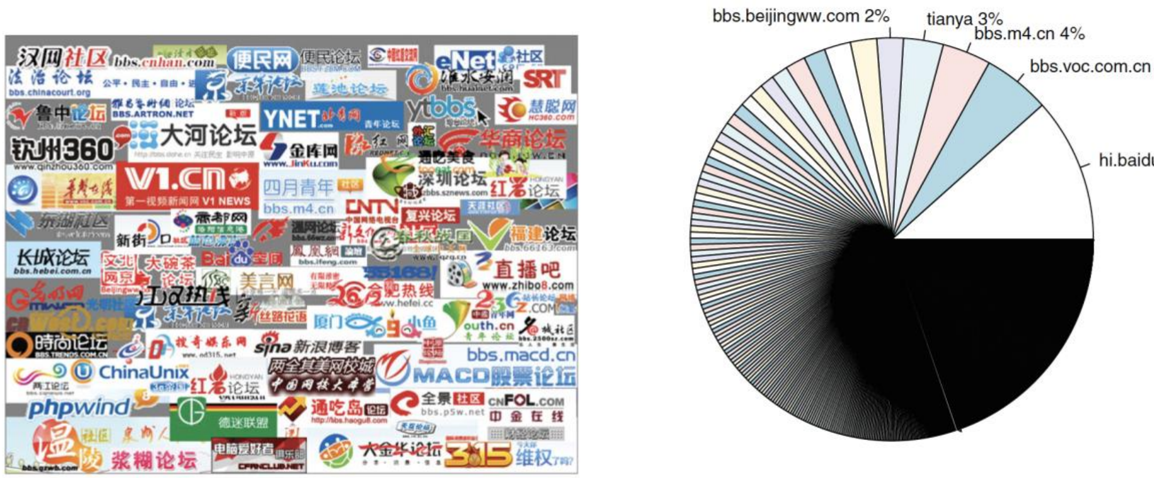

(a) Sample of Sites (b) All Sites excluding Sina

All tables and figures appear in color in the online version. This version can be found at http://j.mp/LdVXqN.

Figure 2. The Speed of Censorship, Monitored in Real-Time

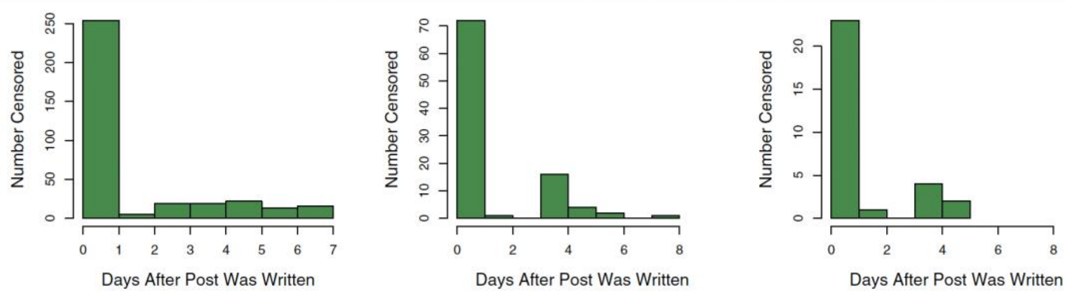

(a) Shanghai Subway Crash (b) Bo Xilai (c) Gu Kailai

Figure 3. "Censorship Magnitude," The Percent of Posts Censored Inside a Volume Burst Minus Outside Volume Bursts.

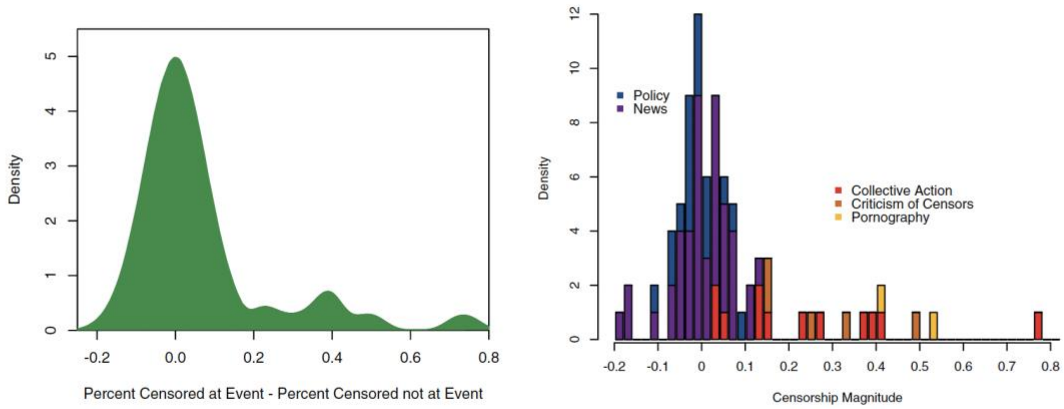

(a) Distribution of Censorship Magnitude

(b) Censorship Magnitude by Event Type

Figure 4. Events with Highest and Lowest Censorship Magnitude

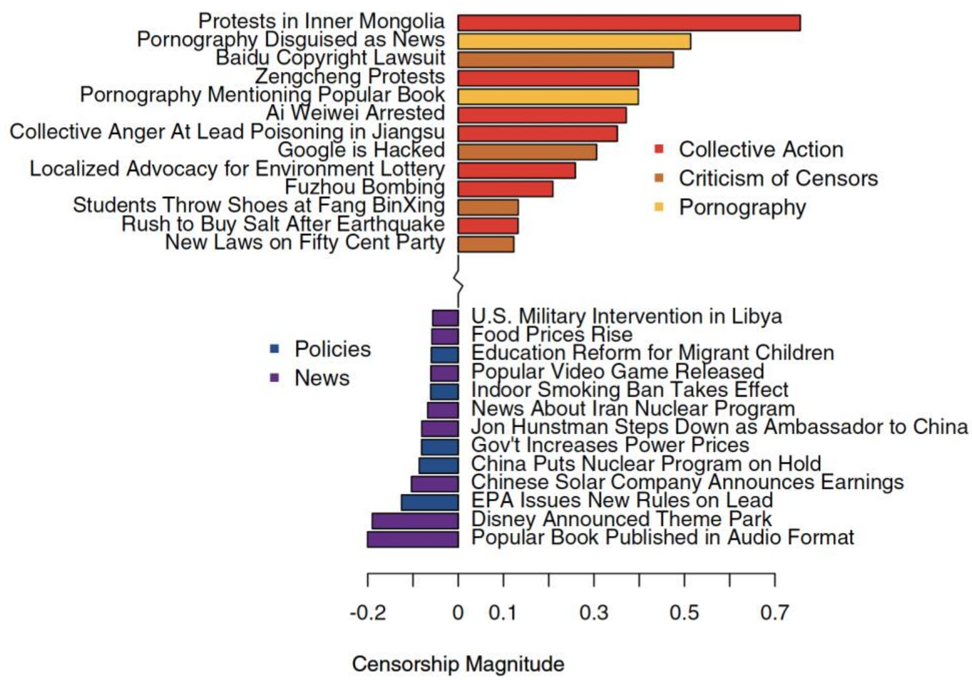

Figure 7. Two Topics with Continuous High Censorship Levels (in 2011)

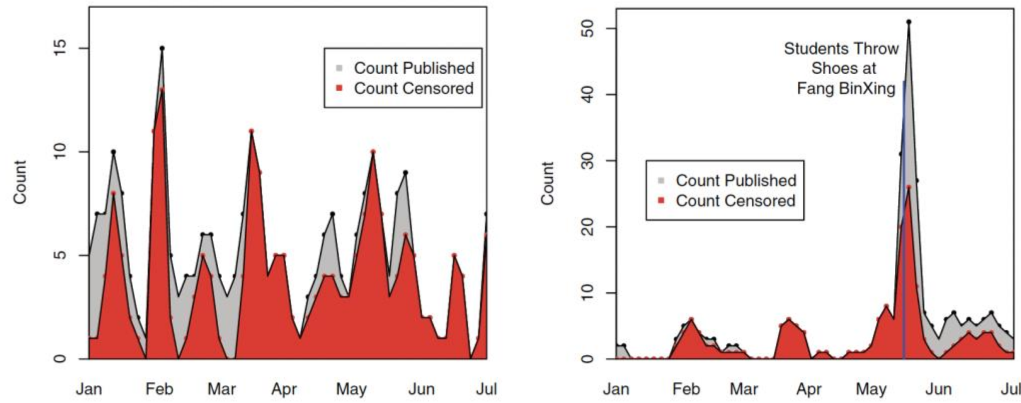

(a) Pornography (b) Criticism of the Censors

c. Authors' Conclusions

The new data and methods we offer seem to reveal highly detailed information about the interests of the Chinese people, the Chinese censorship program, and the Chinese government over time and within different issue areas. These results also shed light on an enormously secretive government program designed to suppress information, as well as on the interests, intentions, and goals of the Chinese leadership.
… leadership… allowed the full range of expression of negative and positive comments about the state, its policies, and its leaders… so long as they manage to eliminate discussions associated with events that have collective action potential.
… the Chinese people are individually free but collectively in chains.

4. Discussion

a. Implications

- The findings suggest that state critique theory does not consistently explain censorship behavior in authoritarian systems. Instead, the theory of collective action potential offers a more appropriate theoretical lens for understanding online censorship in China.
- Censorship pattern can be used as a data source to infer government intent.
- Difference in localized vs. national censorship
- Automated, computational method is able to identify manual censorship effort at large-scale
- Captured and interpreted time-series of events.
- Discovered communities form around new subjects over time.
- Adapted Hopkins & King (2010) in Chinese language for automated text analysis.

b. Limitations

- Removal or deletion of posts didn't consider self-censorship
- Didn't consider the consequences for the users whose posts were censored, especially in highly politically sensitive events.

c. Future Directions

- Can use same method to predict future actions outside the Internet
- Study design can be applied to other topic areas and fields.

d. …

--- Page Break Here ---
Response Letter (~600 words)

- Your assessment of the study
- Weaknesses
- Other notes

This study provides one of the most influential and methodologically innovative analyses of digital censorship in authoritarian contexts. This study investigated censorship in China and found that posts that signal collective action potential are more likely to be censored, especially during politically sensitive events, which is in line with the theory proposed by the authors. At the same time, criticism toward the state did not always end up being censored, which runs contrary to the dominant state critique theory. At the time of publication in 2013, this study pioneered the development of a methodological approach for studying censorship at large scale, offering important methodological advancement in the communication field.

One of the significant strengths of this study is its data collection strategy. The authors took advantage of the fact that censorship in China is largely manual and involves a time lag between original posting and deletion. By implementing an automated system that could identify posts and monitor their removal in real time, the authors were able to observe censorship directly rather than relying on indirect reports or user self-reports in smaller quantities. This approach demonstrated a technical advancement, especially in a highly restricted political environment where access to data is highly limited.

Methodologically, the study is also robust in its integration of time-series analysis, automated text classification, and human coding. Through rigorous human coding, detailed contextualization of local and national events, and automated textual analysis adapted from Hopkins and King (2010), the authors were able to capture large volumes of data and provide empirical evidence for censorship intent. The contextual background of many significant events was particularly helpful in interpreting the patterns and volume bursts observed in the data.

Although the study did not use advanced quantitative modeling or statistical hypothesis testing, the counting and comparison of posts revealed critical insights into the topic they aimed to study, which is censorship in China. Instead of testing multiple statistical hypotheses, the authors relied on a general central hypothesis to evaluate competing theories of censorship. In this case, descriptive analysis was sufficient to reveal systematic patterns in censorship behavior, and the absence of complex models does not weaken the core findings.

Substantively, the findings challenge one of the most common assumptions when it comes to state censorship. The idea that censorship primarily suppresses criticism is assumed in both scholarly and public discourse. However, this study demonstrates that critical content is often tolerated, as long as it does not facilitate collective action. This insight contributes to a more nuanced understanding of how authoritarian regimes managed public expression at least a decade ago.

I do not find major weaknesses in this study, which provided appendices and was honest about not being transparent regarding the tools and mechanisms used for crawling for apparent reasons. As mentioned earlier, some may criticize the simplicity of the analysis, such as relying on counting, but I think that for the purpose of the study it did an adequate job in answering the research questions. Working with hundreds of thousands of posts in an automated manner may raise questions about accuracy, and the training dataset might contain errors that could draw criticism. However, I believe these concerns can be justified, as the authors provided explanations for validity and emphasized the novelty of the study at the time.

On a more personal note, I found it particularly interesting that the death of an Inner Mongolian herder caused by a coal truck driver had one of the highest censorship magnitudes in the study. As a Mongolian, I vaguely remember this story from many years ago, and seeing it analyzed in this academic context made the findings digestible and memorable. Also, the detailed descriptions of real-world events helped connect the computational patterns to online censorship behaviors in China.